

Agenda

- views (last class)
- Intro to Indexing
- How Indexes work?
- Indexes and ranged queries
- cons of indexing
- Indexing on multiple col^m
- Indexing on strings
- LiZ BIT Practical.

start @ 9:10 PM.

Intro to Indexing

"select * from table where id = 'someth'"

for row

$O(n)$?

SLOW !!

(2M) we don't want to iterate over every row all the time.

Library to get a book.

history of India

librarian → pointer

(history section).


<u>us</u>
<u>uk</u>
<u>India</u>

× 30 books

× 30 books

{ India }, 20 books.

(2000) ⇒ (20)



⇒ what is at the
starting page!?

(Table of contents).

⇒ end

(index)

↪ moral : page 34, 39.
sorted

→ Tables are actually stored on the
disk

CPU

L1, L2 caches

fastest.
(small).

Memory (RAM)

(fast).

disk

(slow).

operating system / kernel.

can only read from the memory.







disk reads are slow??

→ avoid reading from disk.

→ slow.

→ expensive





disk

Page 1 | id=1; Ayush

id=10.

Page 999.

id=9990

id(1) is in Page 1 on 0^m row.

id = 789 10³

999 pages. { full table scan }.

"select * from students where id=789"

used PK-index
To become

faster.

Indices essentially help you reduce the look up on the disk for every query.

How do indices work??

→ "select * from table where id = 100"
find page where id (100) is stored.

map

id	page
100	1
102	2
104	4
⋮	

$O(n)$ \nearrow
 $O(1)$ read

without map ($O(n)$)
with map ($O(1)$)

→ "select * from users where name = 'Ayush';"

name page

ayush	[1, 3, 5, 7]
-------	--------------

page 1

ayush

Liv	(3, 4, 5)
sheyas	(6, 10, 11)

page 63

Ayush.

without map (1000) nm pages.

page 253

with a map (3, 4 pages). Ayush.

Ranged queries

"select * from students
where

PSP between 40.1 and 90.1;"

worst case map \Rightarrow range $O(n)$.



- ⚡ ① sorting helps searching Maps are not
- ⚡ ② we want faster look-ups. sorted.



ordered maps

Tree maps

(Balanced BST).



① go to 40.1

② keep going / traversing till next node
> 90.1


Data structure : B Tree / B+ Trees

- ① Are we storing more info in reduced height
- ② can we fetch more info on a single look up?

- ③ B trees are specifically designed for storing data on disks where on a single fetch, we can get more info.

Break Till 10:55

cons of Indexing



Reads are faster!!

- ① writes are slower.
→ also update index

index are also stored on disk.

Indexes are stored in memory.

- ② More storage space on disk.

→ never create indexes prematurely

2×10^9 bytes

Indexing on multiple col^m

id name email PSP.

→ select * from stu where PSP = 80;
index (name, PSP).

① sorted

② lookups

order by (name, PSP)

index (name, PSP).

↑
create
Tree

→ if name match
sort on PSP.

Ayush	30.
Siv	90.
Kanya	60
Surajit	30
Ayush	60.

(PSP, name)	
30	Ayush
90	Siv
60	Ayush.

Index (name, PSP)

≠

Index (name)

&

Index (PSP)

Query on	Index	help	not help.
name	PSP.		x.
name	<u>(name)</u> and (PSP)	✓	
name	(PSP, <u>name</u>)		x.
name.	(name, PSP)	✓	
<u>name = x</u> and <u>PSP = y</u>	(PSP, name)	✓	

name = x and PSP = y

name = x and PSP = y.

name = x


or

PSP = y

(PSP, name).

select * from

where
name = x
PS = 4.



if a query is a col^m x.
any index with x as prefix will
help.

(x, a, b, c) → help.

(x) = help.

(a, x) → won't help.

(x, a) → help.

Indexing on strings

select * from students where
email = "abc@sealar.com"

slow (no index
string matching).

→ should we index the string completely
or partly?

@salar.com.

space

Ayush.saraswat.

Ayush saraswat.

faster??

whole string

email	page
Ayush.saraswat.	1.
Ayush saraswat.	2.

↓
less space

first 5 letter

Ayush

[1, 2]

users 2 Billions

index on number of characters	distinct keys in <u>index</u>	page access/ fetch.
1	26 -	$2B/26$
2	<u>26×26</u>	<u>$2B/(26)^2$</u>
7 char	<u>$(26)^7$</u>	<u>≈ 1</u>

2B

a ... $[2B/26]$

b $[2B/26]$

c

0

7.

increase in number letter

space ↑

fetch ↓

Select * from Table
where address like '% Ambala%'

↑
queries like these, don't use index
and do a full table scan

'ambala%'

could use an index