

## Agenda

- Isolation levels.
- read uncommitted
- read committed
- repeatable reads
- serializable
- dead locks.

start @ 9:10


## Isolation levels

The form in which transactions will be isolated from each other in a database.

- |   |                     |   |                      |
|---|---------------------|---|----------------------|
| ① | Read uncommitted    | → | most relaxed         |
| ② | Read committed      | } | <u>severity</u>      |
| ③ | Repeatable reads    |   |                      |
| ④ | <u>Serializable</u> |   |                      |
|   |                     | ↓ | → <u>most strict</u> |
|   |                     | ↓ | default              |


↓  
generally  
default.

each connections/sessions has its own  
isolation and are also not affected  
by others.



Read uncommitted

Allows a transaction to read uncommitted  
data from another transaction.



$\text{read}(x)$   
 $?? = (1)$   
 $(0, 1)$

$\text{update}(x)$   
 $= 0$

→ basically means that you will read the latest data.

less strict ↓

fast ↑

A = 2000  
B = 2000

A = 2100  
B = 1900

A  $\xrightarrow{10/-}$  B ✗  
 Transaction 1 (RR)

B  $\xrightarrow{100/-}$  A  
 Transaction 2 (RU)

- ① Read A →  $x(2000)$
- ②  $x \rightarrow x - 10$   $VA = 1990$   
 $CA = 2000$   
 $CB = 2000$
- ③ write A ←  $x(1990)$

- ① Read B →  $x(2000)$
- ②  $x \rightarrow x - 100$

- ④ Read B →  $x(2000)$

- ③ write B ←  $x(1900)$

- ⑧ Read A →  $x(1990)$

fails  
 & roll  
 Backs

- ⑨  $x \rightarrow x + 100(2090)$

X. dirty read

⑩ write to A  
 $A \leftarrow x(2090)$   
commit;

100%  
A = 2090  
B = 1900

→ T<sub>1</sub> rolled back; T<sub>2</sub> (RC) continued

→ dirty read: the state of reading uncommitted data.


→ we are not sure about the consistency of data that is read because we don't know the result of a an open transact<sup>n</sup>.

→ the open T<sub>1</sub> MIGHT commit or MIGHT rollback.

Read committed.

↑ data.

reads. → allowed.  
writes

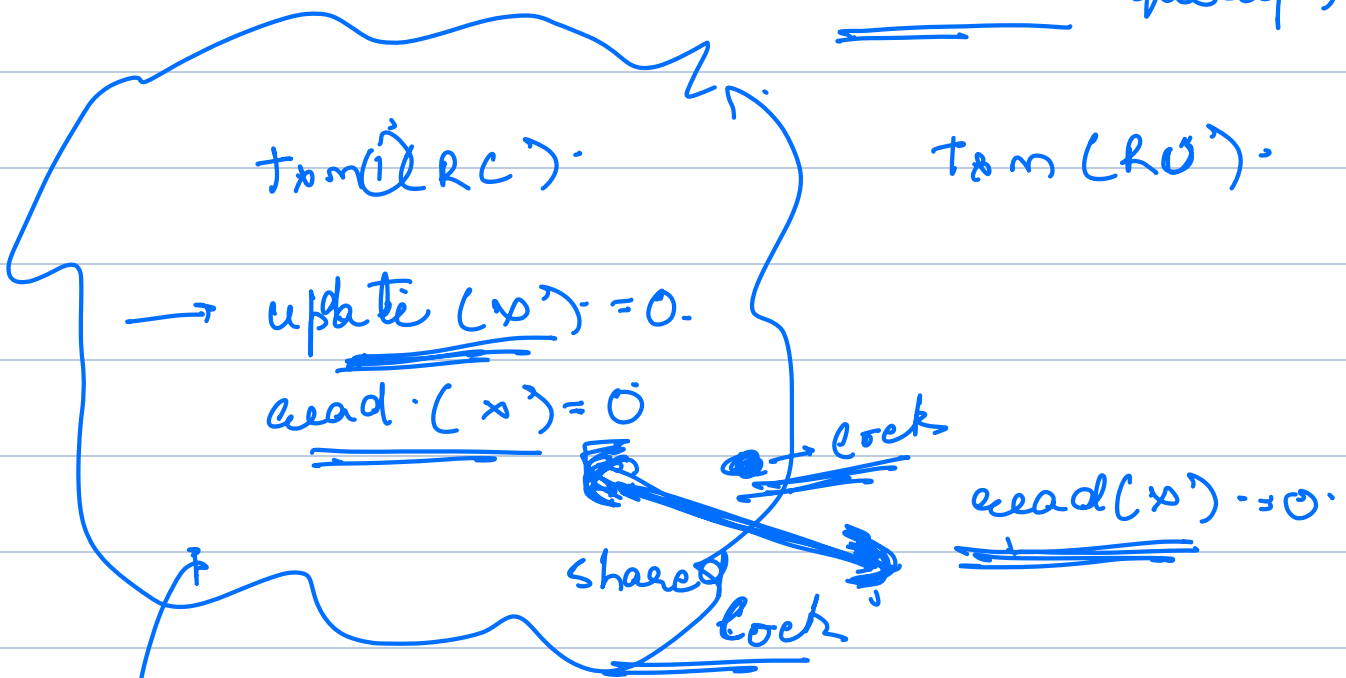


from diff txn

→ update on same rows.  
you get a halt. ~

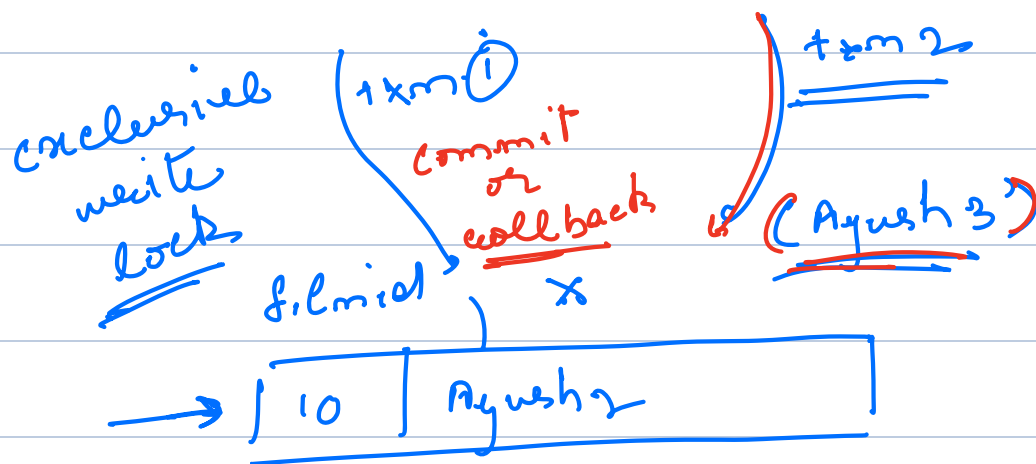
→ Becz the 1st Transact<sup>m</sup> aquired  
a Lock.

→ shared lock: (caused by a read  
query).



~~txn (3) (RR)~~  
 read (n)  
disk

→ exclusive lock (updates / write query).



→ only and only one transact<sup>m</sup> can write on a row when concurrent txns are going on.

Read committed :

- ① read values which are committed.
- ② No dirty reads.

RC is slower than RV.

→ validation

→ lots of concurrent control  
mechanisms.

break till 10:42

Problems with read committed

id	name	psp	is_email_sent?
1	m	60	false
2	A	(80)	false
3	m	70	false
4	ab	45	false.
5	ac	35	

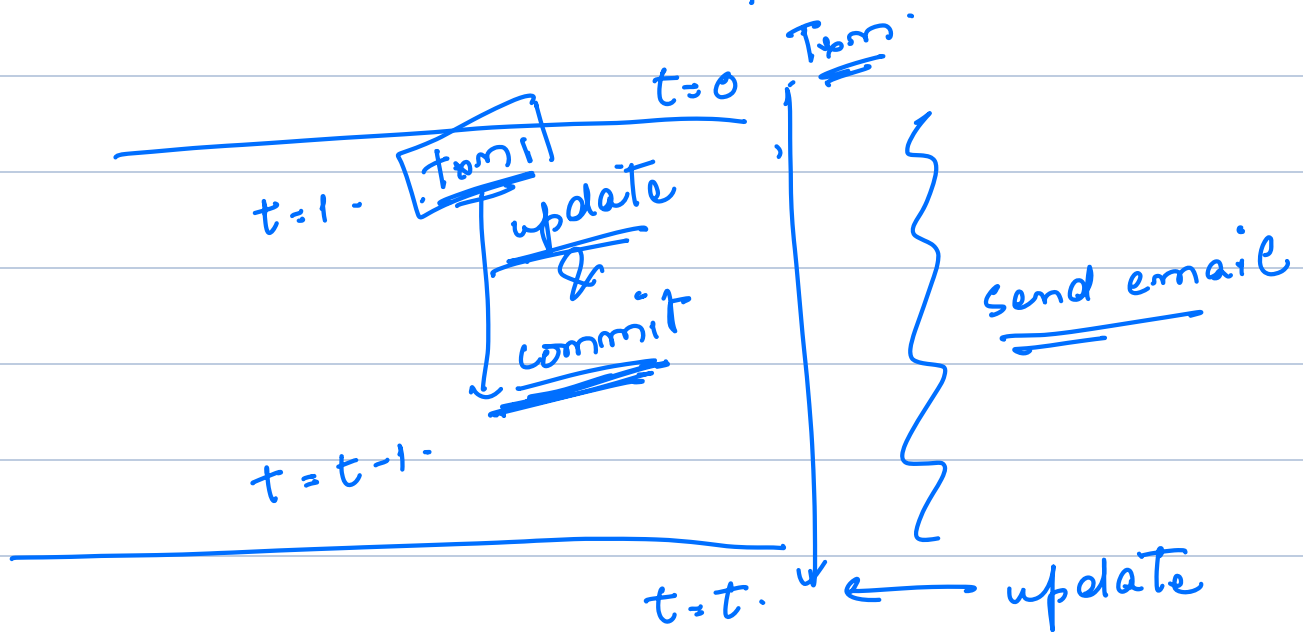
→ I want to send an email  
to students psp < 80

step 1: getting list of users  $psp < 80$   
(id 2 will definitely)  
— not be there.

step 2: sending emails to users  
in the list.

(2 will not get the mail)

step 3: due to some reason the PSP of user 2  
was changed to 79.



step (4): while updating "is\_email\_sent"  
we go through all the values  
with PSP  $< 80$

update students

set is\_email\_sent = true

where PSP  $< 80$

80 79 ?

Did we send an email to user 2?




NO

was it marked as sent ??

Yes!!

with RR  
80, 80, 80  
80, 80

NO!!



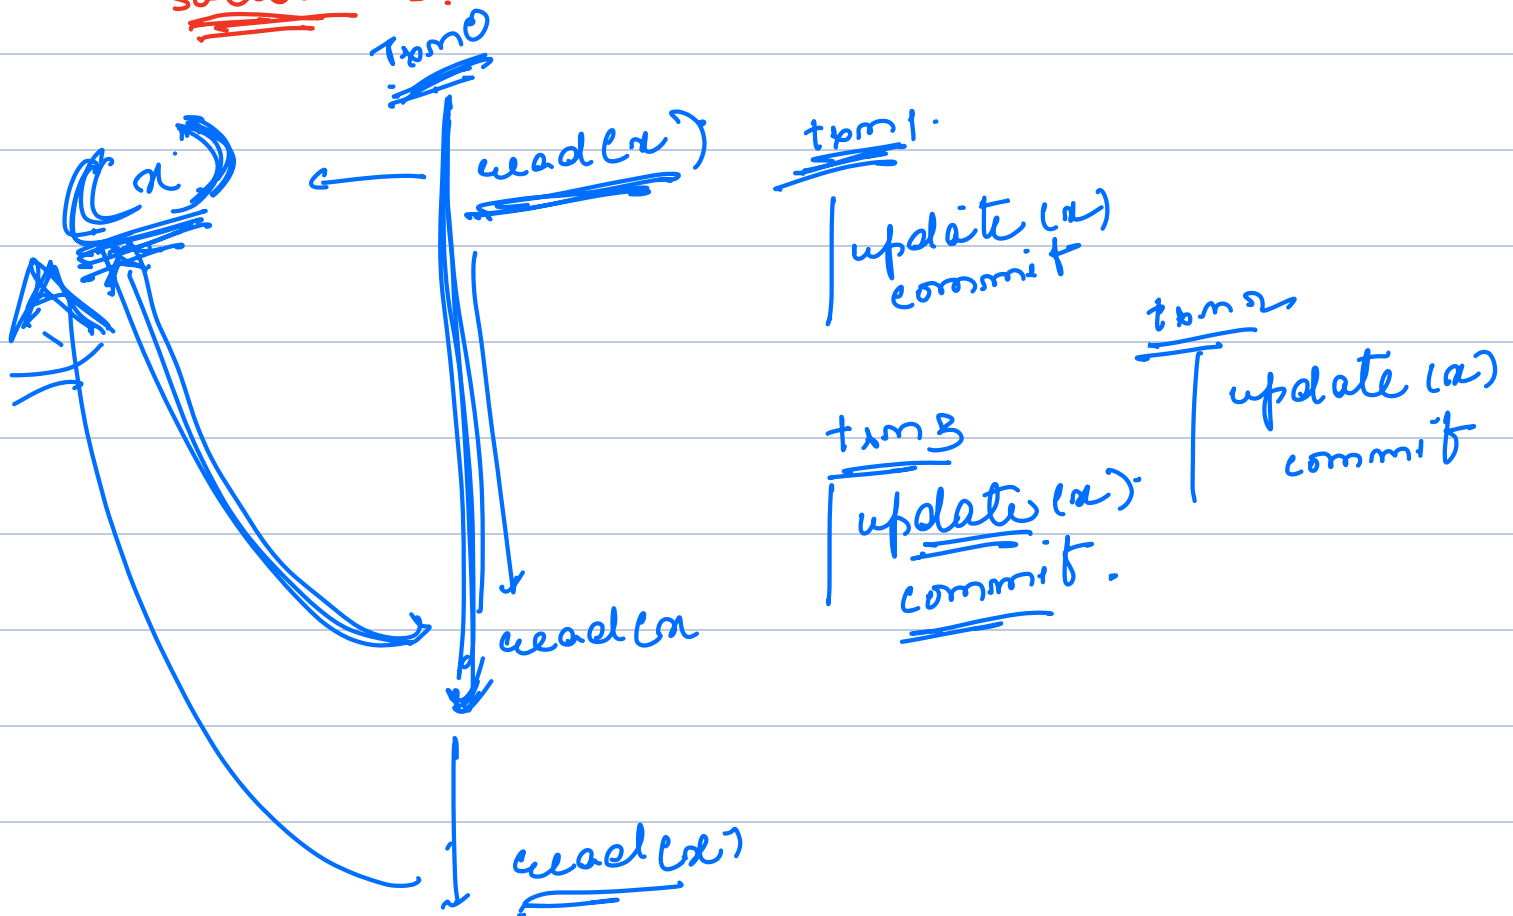
expectat<sup>n</sup>  
same.

reality  
not same.

non - Repeatable read:

a transact<sup>n</sup> makes 2 reads  
at diff times and get diff  
values.

Solution:



## Repeatable reads

- for the first time, it reads the latest committed.
- after that, until the tm completes, it keeps on reading the same value as the first time.

Dirty reads ? ? yes

Non repeatable → yes Repeatable  
read isolation

check  
→ if we read a row again, will we get the same value.

→ 100 → 120


100  
< 80 Pst & <

Phantom reads:

the value of repeatable reads are same, but more or less number of rows satisfy your read conditions.

100 rows → 120 rows.  
(agg) - (agg1) -

agg = agg1 ✗



PSP  
280

↓ read (102)  
↓ agg (sum(PSP))

commit;